


Learned dynamic K uses extra refinement sparsely


Distribution of selected recurrent depth over all imagined CEM transitions

K=1 K=2 K=3 K=4

All selected depths



Zoom: transitions refined beyond K=1



Takeaway: the learned selector improves success while sending only a small fraction of transitions past K=1.