


Where does learned dynamic K spend extra refinement?



Each point aggregates all candidate rollouts and MPC calls in the evaluation. Higher values mean more transitions continued to $K>1$ at that rollout step.